

```

1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 import warnings
6 %matplotlib inline
7 from scipy.stats import normaltest, shapiro, anderson
8
9 warnings.filterwarnings('ignore')

```

```
1 df = pd.read_csv('/content/gas_daily.csv')
```

```
1 df.head(10)
```

	Date	Price	
0	07/01/1997	3.82	
1	08/01/1997	3.80	
2	09/01/1997	3.61	
3	10/01/1997	3.92	
4	13/01/1997	4.00	
5	14/01/1997	4.01	
6	15/01/1997	4.34	
7	16/01/1997	4.71	
8	17/01/1997	3.91	
9	20/01/1997	3.26	

```
1 df.describe()
```

	Price	
count	5952.000000	
mean	4.184644	
std	2.190361	
min	1.050000	
25%	2.650000	
50%	3.530000	
75%	5.240000	
max	18.480000	

```
1 df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5953 entries, 0 to 5952
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Date    5953 non-null    object
1    Price    5952 non-null    float64
dtypes: float64(1), object(1)
memory usage: 93.1+ KB

```

```
1 df.shape
```

```
(5953, 2)
```

```
1 df.isna().sum()
```

```

Date      0
Price     1
dtype: int64

```

```
1 df.dropna(inplace=True)
```

```
2
```

```
1 df['Date'].max()

'31/12/2019'


1 df['Date'].min()

'01/01/2016'
```

LINE CHART

```
1 df['Date'] = pd.to_datetime(df['Date'])

1 plt.style.use('seaborn-darkgrid')
2
3 plt.figure(figsize=(12, 6))
4 plt.plot(df['Date'], df['Price'], label='Natural Gas Price', color='pink')
5 plt.title('Natural Gas Price Over the Last 24 Years')
6 plt.xlabel('Date')
7 plt.ylabel('Price')
8 plt.legend()
9 plt.grid(True)
10 plt.show()
11
12 plt.figure(figsize=(12, 6))
13 plt.fill_between(df['Date'], df['Price'], color='pink', alpha=0.2, label='Natural Gas Price')
14 plt.title('Area Chart: Natural Gas Price Over the Last 24 Years')
15 plt.xlabel('Date')
16 plt.ylabel('Price')
17 plt.legend()
18 plt.grid(True)
19 plt.show()
```



The line chart indicates a distinct trend in gas prices over the last two decades. Following a post-2000 upturn, prices remained consistently high between 2000 and 2009. However, a subsequent decline is evident, marking a shift in market dynamics.

Explanatory variables

```
1 df.sort_values(by='Date', inplace=True)
2
3 df['MA5'] = df['Price'].rolling(window=5).mean()
4 df['MA10'] = df['Price'].rolling(window=10).mean()
5 df['MA20'] = df['Price'].rolling(window=20).mean()
6
7 df.dropna(inplace=True)
8
9 print(df[['Date', 'Price', 'MA5', 'MA10', 'MA20']])
```

	Date	Price	MA5	MA10	MA20
17	1997-01-30	2.86	2.884	3.059	3.0895
18	1997-01-31	2.77	2.914	2.945	3.1305
59	1997-02-04	1.85	2.688	2.804	3.1135
81	1997-02-05	2.21	2.520	2.726	3.1160
101	1997-02-06	2.20	2.378	2.641	3.1140
...
5808	2020-12-02	1.91	1.906	1.898	1.9625
5830	2020-12-03	1.82	1.878	1.889	1.9275
5875	2020-12-05	1.61	1.860	1.878	1.8850
5897	2020-12-06	1.67	1.840	1.866	1.8535
5938	2020-12-08	2.05	1.812	1.853	1.8535

[5933 rows x 5 columns]

Training & Testing

```
1 from sklearn.model_selection import train_test_split
2
3
4 features = df[['MA5', 'MA10', 'MA20']]
5 target = df['Price']
6
7 X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)
8
9 print("X_train shape:", X_train.shape)
10 print("X_test shape:", X_test.shape)
11 print("y_train shape:", y_train.shape)
12 print("y_test shape:", y_test.shape)
```

```
X_train shape: (4746, 3)
X_test shape: (1187, 3)
y_train shape: (4746,)
y_test shape: (1187,)
```

```
1 features = df[['MA5', 'MA10', 'MA20']]
2 target = df['Price']
```

```
1 X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)
```

```
1 print("X_train shape:", X_train.shape)
2 print("X_test shape:", X_test.shape)
3 print("y_train shape:", y_train.shape)
4 print("y_test shape:", y_test.shape)
```

```
X_train shape: (4746, 3)
X_test shape: (1187, 3)
y_train shape: (4746,)
y_test shape: (1187,)
```

Building Linear Regression Model

```

1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import mean_squared_error, r2_score
3 import numpy as np
4
5
6 model = LinearRegression()
7
8 model.fit(X_train, y_train)
9
10 train_predictions = model.predict(X_train)
11 test_predictions = model.predict(X_test)
12
13 train_rmse = np.sqrt(mean_squared_error(y_train, train_predictions))
14 test_rmse = np.sqrt(mean_squared_error(y_test, test_predictions))
15 r2_train = r2_score(y_train, train_predictions)
16 r2_test = r2_score(y_test, test_predictions)
17
18 print("Training RMSE:", train_rmse)
19 print("Testing RMSE:", test_rmse)
20 print("Training R^2 Score:", r2_train)
21 print("Testing R^2 Score:", r2_test)

```

```

Training RMSE: 0.7023196191348168
Testing RMSE: 0.6744025354238129
Training R^2 Score: 0.8973598880789414
Testing R^2 Score: 0.9053319409388314


```

Prediction Function and Result

```

1 import matplotlib.pyplot as plt
2 plt.figure(figsize=(12, 6))
3 plt.plot(X_test.index, y_test, label='Actual Prices', color='blue')
4 plt.plot(X_test.index, test_predictions, label='Predicted Prices', color='red', linestyle='dashed')
5 plt.title('Actual vs. Predicted Natural Gas Prices')
6 plt.xlabel('Date')
7 plt.ylabel('Price')
8 plt.legend()
9 plt.grid(True)
10 plt.show()
11
12
13 r_squared = model.score(X_test, y_test)
14 accuracy = r_squared * 100
15 print("Model Accuracy (R^2 Score):", accuracy)

```



Model Accuracy (R² Score): 90.53319409388314

alpha and betas value

```

1 alpha = model.intercept_
2 betas = model.coef_
3 print(alpha, betas)

```

```
5 print(alpha,betas)
```

```
0.018793686576505486 [ 1.18598754 -0.53534093  0.34660581]
```

```
1 alpha = model.intercept_
```

```
2 betas = model.coef_
```

```
3
```

```
4 features = ['MA5', 'MA10', 'MA20']
```

```
5
```

```
6 equation = f"Y = {alpha:.2f} + " + " + ".join([f"{beta:.2f} * {feature}" for beta, feature in zip(betas, features)])
```

```
7 print("Linear Regression Equation:")
```

```
8 print(equation)
```

```
Linear Regression Equation:
```

```
Y = 0.02 + 1.19 * MA5 + -0.54 * MA10 + 0.35 * MA20
```